



Ritik Kumar

DATA ENGINEER

+91 8271448465

ritik74820@gmail.com

PORTFOLIO QR



Summary

Data Engineer with a strong interest in AI and Machine Learning. Focused on building reliable data pipelines, transforming and cleaning data, and turning raw data into structured, usable formats for analytics and intelligent applications. Works with SQL, Python, ETL/ELT pipelines, data modeling, databases, and cloud-based data engineering. Currently deepening knowledge of AI/ML, with the goal of becoming a strong Data Engineer capable of handling large-scale data systems.

Skills

Languages:

- Python
- SQL
- T-SQL
- Bash

Data Processing:

- PySpark
- Pandas
- Polars

Data Engineering:

- Data Modeling
- Data Warehousing
- ETL / ELT
- Data Pipeline
- Database Architecture

Artificial Intelligence:

- Prompting
- Generative AI
- AI CLI Tools

Certifications

- Data Engineer Certification
- Data Engineer Associate
- Big Data with PySpark
- SQL Associate
- Data Engineering Hands-On Practice
- Data Processing and Pipeline Optimization

Education

Higher Secondary (12th)

Bihar School Examination Board (BSEB)

Subjects: PCM | English Medium

Year: 2024

Languages

- English Full Working Proficiency
- Hindi Native Proficiency
- Maithili / Magahi Native Proficiency

Few Other Indian Language

Personal Interests

- Book Reading & Self-Learning
- Fitness, Yoga & Meditation
- Discipline & Consistent Routines
- Entrepreneurship

Project Experience

Retail Analytics Engineering Platform

Tech: dbt Core • SQL Server 2022 • Docker • GitHub Actions

Engineered a dbt-based Medallion pipeline on SQL Server, transforming retail data into a Kimball Star Schema with custom macros, SCD Type 2 snapshots, data-quality tests, and GitHub Actions CI/CD.

SkyNova Data Ecosystem

Tech : Snowflake • dbt • Python • SQL • AWS • ML • AI

Designed a synthetic dirty-data specification for a 9-source, 1.9M-row retail dataset, modeling SCD2 history, identity resolution, and 8+ data-quality defect classes for a Kimball dbt/Snowflake pipeline.

Medallion Data Warehouse

Tech: Microsoft SQL Server • T-SQL • Docker

Built a three-layer T-SQL data warehouse for 8 retail domains, implementing defensive data cleansing, standardization, and a star schema with 7 dimensions and 3 fact views for sales, inventory, and returns.

Snowflake Semi-Structured Data

Tech: Snowflake • SQL • VARIANT/JSON • Python

Implemented Snowflake semi-structured data pipelines using VARIANT, JSON parsing, and LATERAL FLATTEN across deeply nested data, normalizing it into a 3NF data mart with idempotent MERGE-based incremental loading.

Great Minds Knowledge Graph

Tech: React • TypeScript • PostgreSQL • Kysely • AI-Assisted

Built an AI-assisted research knowledge graph application using agentic AI and vibe coding, modeling people, decisions, failures, evidence, confidence, and source provenance into an interactive research platform.